

# SAMformer: Unlocking the Potential of Transformers in Time Series Forecasting

Ambroise Odonnat

Huawei Noah's Ark Lab

Criteo AI Lab

July 15, 2024



**CRITEO**  
**AI Lab**




- ① Introduction
- ② Failure of Transformers
- ③ SAMformer
- ④ Experiments
- ⑤ Take Home Message



- ① Introduction
- ② Failure of Transformers
- ③ SAMformer
- ④ Experiments
- ⑤ Take Home Message

In many applications, data are gathered sequentially.





Goal → Analysing time series data to predict future trends.

- Forecast of ECG recording to predict cardiac arrhythmia,
- Electricity consumption forecasting to match future demand,
- Predicting stock market prices.

## Challenges



- ① Long-term temporal dependencies,
- ② Highly correlated features.



$D$ -dimensional time series of length  $L \rightarrow$  predict next  $H$  values.

- Training set of  $N$  observations  $(\{\mathbf{X}^{(i)}\}_{i=0}^N, \{\mathbf{Y}^{(i)}\}_{i=0}^N)$ ,
- Find predictor  $f_{\omega}: \mathbb{R}^{D \times L} \rightarrow \mathbb{R}^{D \times H}$  that minimizes the MSE

$$\mathcal{L}_{\text{train}}(\omega) = \frac{1}{ND} \sum_{i=0}^N \|\mathbf{Y}^{(i)} - f_{\omega}(\mathbf{X}^{(i)})\|_F^2.$$








$D$ -dimensional time series of length  $L \rightarrow$  predict next  $H$  values.

- Training set of  $N$  observations  $(\{\mathbf{X}^{(i)}\}_{i=0}^N, \{\mathbf{Y}^{(i)}\}_{i=0}^N)$ ,
- Find predictor  $f_{\omega}: \mathbb{R}^{D \times L} \rightarrow \mathbb{R}^{D \times H}$  that minimizes the MSE

$$\mathcal{L}_{\text{train}}(\omega) = \frac{1}{ND} \sum_{i=0}^N \|\mathbf{Y}^{(i)} - f_{\omega}(\mathbf{X}^{(i)})\|_F^2.$$






$D$ -dimensional time series of length  $L \rightarrow$  predict next  $H$  values.

- Training set of  $N$  observations  $(\{\mathbf{X}^{(i)}\}_{i=0}^N, \{\mathbf{Y}^{(i)}\}_{i=0}^N)$ ,
- Find predictor  $f_{\omega}: \mathbb{R}^{D \times L} \rightarrow \mathbb{R}^{D \times H}$  that minimizes the MSE

$$\mathcal{L}_{\text{train}}(\omega) = \frac{1}{ND} \sum_{i=0}^N \|\mathbf{Y}^{(i)} - f_{\omega}(\mathbf{X}^{(i)})\|_{\text{F}}^2.$$







## Standard methods:

- AR models (ARIMA)
- Seasonal naive

## Deep learning methods

- RNN, CNN
- Transformer-based models



- ① Introduction
- ② Failure of Transformers
- ③ SAMformer
- ④ Experiments
- ⑤ Take Home Message



## Motivation

- Transformers tailored to deal with sequential data,
- Impressive results in NLP and Computer Vision.

## Main challenges

- ① Quadratic computation of self-attention,
- ② Complex long-term dependencies.



## Main challenges

- ① Quadratic computation of self-attention,
- ② Complex long-term dependencies.



## Main challenges

- ① Quadratic complexity of self-attention
  - Sparse attention: LogTrans [5], Informer [13]
  - Modified attention: Pyraformer [6]
- ② Complex long-term dependencies
  - Decomposition scheme: Autoformer [10], Pyraformer [6]
  - Fourier domain: FEDformer [14]

**It leads to a wide range of Anything-formers with heavy and complex implementation and many parameters.**

[11] showed that linear models outperform SOTA Anything-former.

## Transformers in Computer Vision and NLP



**Commander of the Armies of GPT,  
General of the Gemini Legions, loyal  
servant to Claude, Llama3, Mixtral**

## Transformers in Time Series Forecasting




**Please help, I just got  
beaten by a linear model**



- ① Introduction
- ② Failure of Transformers
- ③ SAMformer
- ④ Experiments
- ⑤ Take Home Message




- Generate toy data according to  $\mathbf{Y} = \mathbf{X}\mathbf{W}_{\text{toy}} + \epsilon$ ,
- Designing the simplest Transformer possible.








- Generate toy data according to  $\mathbf{Y} = \mathbf{XW}_{\text{toy}} + \epsilon$ ,
- Designing the simplest Transformer possible.





- Generate toy data according to  $\mathbf{Y} = \mathbf{X}\mathbf{W}_{\text{toy}} + \epsilon$ ,
- Designing the simplest Transformer possible.


$$\mathbf{A}(\mathbf{X}) = \text{softmax}\left(\frac{\mathbf{X}\mathbf{W}_Q\mathbf{W}_K^\top\mathbf{X}^\top}{\sqrt{d_m}}\right)$$
$$f(\mathbf{X}) = [\mathbf{X} + \mathbf{A}(\mathbf{X})\mathbf{X}\mathbf{W}_V\mathbf{W}_O]\mathbf{W}$$





- Generate toy data according to  $\mathbf{Y} = \mathbf{X}\mathbf{W}_{\text{toy}} + \epsilon$ ,
- Designing the simplest Transformer possible.

$$\mathbf{A}(\mathbf{X}) = \text{softmax}\left(\frac{\mathbf{X}\mathbf{W}_Q\mathbf{W}_K^\top\mathbf{X}^\top}{\sqrt{d_m}}\right)$$
$$f(\mathbf{X}) = [\mathbf{X} + \mathbf{A}(\mathbf{X})\mathbf{X}\mathbf{W}_V\mathbf{W}_O]\mathbf{W}$$




Theorem (Ilbert, O., Feofanov et al.)

*Given fixed attention weights  $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V, \mathbf{W}_O$ , there exists an infinity of optimal  $\mathbf{W}$  reaching the oracle, i.e.,  $f(\mathbf{X}) = \mathbf{X}\mathbf{W}_{\text{toy}}$ .*



- Generate toy data according to  $\mathbf{Y} = \mathbf{X}\mathbf{W}_{\text{toy}} + \varepsilon$ ,
- Designing the simplest Transformer possible


$$\mathbf{A}(\mathbf{X}) = \text{softmax}\left(\frac{\mathbf{X}\mathbf{W}_Q\mathbf{W}_K^\top\mathbf{X}^\top}{\sqrt{d_m}}\right)$$
$$f(\mathbf{X}) = [\mathbf{X} + \mathbf{A}(\mathbf{X})\mathbf{X}\mathbf{W}_V\mathbf{W}_O]\mathbf{W}$$




**In theory, our simplistic Transformer can be optimal. Is this the case in practice?**



- Oracle: optimal solution,
- Transformer with  $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V, \mathbf{W}_O, \mathbf{W}$  trainable,
- Random Transformer: only  $\mathbf{W}$  is trainable.




Despite its simplicity, Transformer overfits a lot. Fixing the attention weight improves generalization.



Poor generalization of Transformer with SGD, Adam, and AdamW.

## Hypothesis from NLP and Computer Vision

- Transformers have **sharp loss landscape** [2]
  - Convergence to sharp minima  $\rightarrow$  poor generalization,
  - Computed as  $\lambda_{\max}$ , the largest singular value of the Hessian.




## Hypothesis from NLP and Computer Vision

- Transformers have **sharp loss landscape** [2],
- Attention suffers from **entropy collapse** [12].
  - Entropy = average entropy of the rows,
  - It causes training instability,
  - [12] → entropy collapse and sharpness appear in tandem.



~ Uniform → high entropy.




Diagonal → low entropy.



## Hypothesis from NLP and Computer Vision

- Transformers have **sharp loss landscape** [2],
- Attention suffers from **entropy collapse** [12].



Training the attention induces an entropy collapse and a sharp loss landscape.



## ① $\sigma$ Reparam [12]

Replace each weight matrix  $\mathbf{W}$  by

$$\widehat{\mathbf{W}} = \frac{\gamma}{\|\mathbf{W}\|_2} \mathbf{W}, \text{ with } \gamma \in \mathbb{R} \text{ learnable ,}$$

## ② Sharpness-Aware Minimization (SAM) [3]


Replace the training loss  $\mathcal{L}_{\text{train}}$  by

$$\mathcal{L}_{\text{train}}^{\text{SAM}}(\omega) = \max_{\|\epsilon\| < \rho} \mathcal{L}_{\text{train}}(\omega + \epsilon) \approx \mathcal{L}_{\text{train}}\left(\omega + \rho \cdot \frac{\nabla \mathcal{L}_{\text{train}}(\omega)}{\|\nabla \mathcal{L}_{\text{train}}(\omega)\|_2}\right).$$

$\sigma$ Reparam doesn't solve the problem, but SAM does.




Contrary to NLP and Computer Vision, entropy collapse seems benign in time series forecasting while sharpness is harmful.



$\sigma$ Reparam doesn't solve the problem, but SAM does.



Contrary to NLP and Computer Vision, entropy collapse seems benign in time series forecasting while sharpness is harmful.



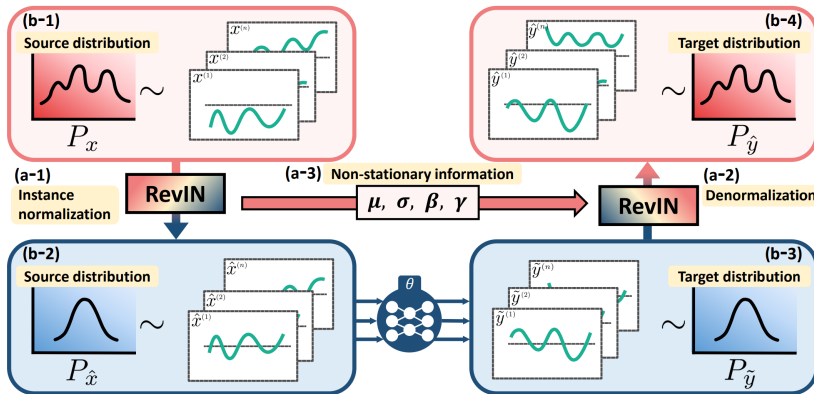
$\sigma$ Reparam helps but is not sufficient while using SAM leads to the optimal solution (oracle).



RevIN [4]  $\rightarrow$  Reduce distribution shift between input and target.

- Multivariate time series  $\mathbf{X} \in \mathbb{R}^{D \times L}$ , learnable  $\gamma, \beta$ .
- Normalize each feature  $\mathbf{X}_i \leftarrow \tilde{\mathbf{X}}_i = \frac{\mathbf{X}_i - \mu}{\sigma} \leftarrow \gamma \tilde{\mathbf{X}}_i + \beta$ ,
- Apply model on multivariate time series  $\tilde{\mathbf{Y}} = f(\tilde{\mathbf{X}})$ ,
- Denormalize each feature  $\tilde{\mathbf{Y}}_i \leftarrow \hat{\mathbf{Y}}_i = \frac{\tilde{\mathbf{Y}}_i - \beta}{\gamma} \leftarrow \sigma \hat{\mathbf{Y}}_i + \mu$ .


RevIN [4] → Reduce distribution shift between input and target.





- Input  $\mathbf{X} \in \mathbb{R}^{D \times L}$ , output  $f(\mathbf{X}) \in \mathbb{R}^{D \times H}$ ,
- Reduce distribution shift with **RevIN** [4],
- **Channel-wise attention**  $\mathbf{A}(\mathbf{X}) \in \mathbb{R}^{D \times D}$ ,
- **Smooth** loss landscape with SAM [3].

$$\mathbf{A}(\mathbf{X}) = \text{softmax}\left(\frac{\mathbf{X}\mathbf{W}_Q\mathbf{W}_K^\top\mathbf{X}^\top}{\sqrt{d_m}}\right)$$
$$f(\mathbf{X}) = [\mathbf{X} + \mathbf{A}(\mathbf{X})\mathbf{X}\mathbf{W}_V\mathbf{W}_O]\mathbf{W}$$



**SAMformer** is a shallow transformer trained with SAM.

→ One head, one encoder, 15 lines of code!



- ① Introduction
- ② Failure of Transformers
- ③ SAMformer
- ④ Experiments
- ⑤ Take Home Message





All-MLP model (2023): TSMixer [1].

Transformers (2021-2022): FEDformer [14], Autoformer [10].

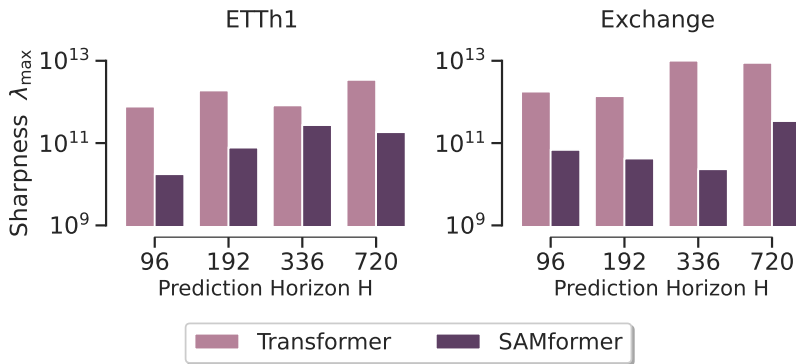
Recent Transformers (2023-2024): iTransformer [7], PatchTST [8].

Dataset	ETTh1/ETTh2	ETTM1/ETTM2	Electricity	Exchange	Traffic	Weather
# features	7	7	321	8	862	21
# time steps	17420	69680	26304	7588	17544	52696
Granularity	1 hour	15 minutes	1 hour	1 day	1 hour	10 minutes

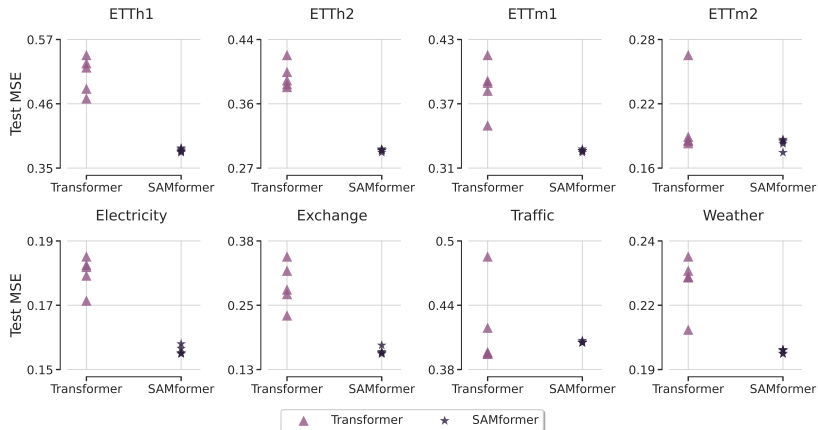


Dataset	<b>SAMformer</b>	iTransformer	PatchTST	TSMixer	FEDformer	Autoformer
	-	2024	2023	2023	2022	2021
ETTh1	<b>0.410</b>	0.454	0.469	0.437	0.440	0.496
ETTh2	<b>0.344</b>	0.383	0.387	0.357	0.437	0.450
ETTm1	<b>0.373</b>	0.407	0.387	0.385	0.448	0.588
ETTm2	<b>0.269</b>	0.288	0.281	0.289	0.305	0.327
Traffic	<b>0.425</b>	0.428	0.481	0.620	0.610	0.628
Weather	0.260	<b>0.258</b>	0.259	0.267	0.309	0.338
<b>Overall improvement</b>		<b>6.58%</b>	<b>8.79%</b>	<b>13.2%</b>	<b>22.5%</b>	<b>35.9%</b>

**SAMformer** outperforms all baselines while having significantly fewer parameters.



**SAM provides a smoother loss landscape ...**



... leading to better generalization and robustness.

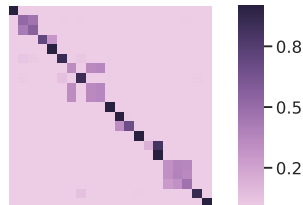
Transformer



$\sigma$ Reparam



SAMformer

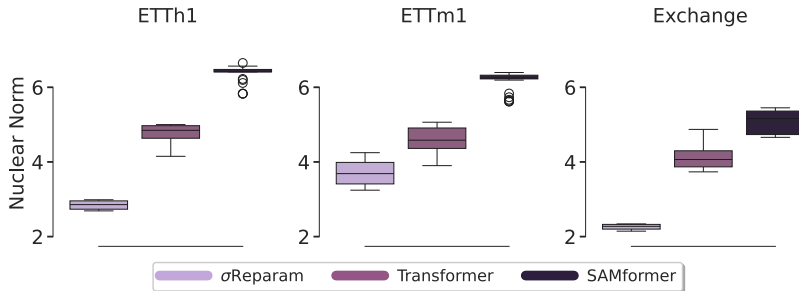


**Channel-wise attention improves the propagation of the signal with self-feature correlations as in ViTs.**

Theorem (Ilbert, O., Feofanov et al.)

Applying  $\sigma$ Reparam [12] leads to **attention rank collapse**.

$$\|\mathbf{X}\mathbf{W}_Q\mathbf{W}_K^\top\mathbf{X}^\top\|_* \leq \underbrace{\|\mathbf{W}_Q\mathbf{W}_K^\top\|_2}_{\text{goes to 0 with } \sigma\text{Reparam}} \|\mathbf{X}\|_F^2.$$





- MOIRAI [9]: foundation model trained on **27B samples**,
- Nb. params: small (**14M**), base (**91M**) and large (**314M**).

Dataset	Full-shot	Zero-shot		
	<b>SAMformer</b>	MOIRAI <sub>Small</sub>	MOIRAI <sub>Base</sub>	MOIRAI <sub>Large</sub>
ETTh1	<u>0.410</u>	<b>0.400</b>	0.434	0.510
ETTh2	<u>0.344</u>	<b>0.341</b>	0.345	0.354
ETTm1	<b>0.373</b>	0.448	<u>0.381</u>	0.390
ETTm2	<b>0.269</b>	0.300	<u>0.272</u>	0.276
Electricity	<b>0.181</b>	0.233	<u>0.188</u>	<u>0.188</u>
Weather	0.260	<u>0.242</u>	<b>0.238</b>	0.259
<b>Overall MSE improvement</b>		<b>6.9%</b>	<b>1.1%</b>	<b>7.6%</b>

**SAMformer outperforms MOIRAI while having significantly fewer parameters!**



- ① Introduction
- ② Failure of Transformers
- ③ SAMformer
- ④ Experiments
- ⑤ Take Home Message





## Findings

- Transformer failure → trainability issues of the attention,
- In time series forecasting, entropy collapse is benign,
- But sharpness prevents good generalization.

## Proposal

- **SAMformer**: RevIN + channel-wise attention + SAM,
- SOTA and lightest model,
- Strong competitor to MOIRAI [9].




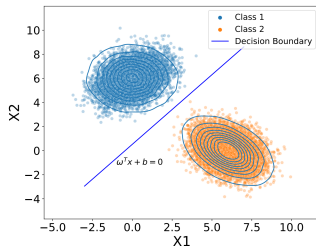
This work has been accepted as an **Oral at ICML 2024, Vienna**. You may find the links to the paper and the code below. To know more about my research, check my website: [ambroiseodt.github.io](https://ambroiseodt.github.io) and feel free to contact me.

- ★ Paper: <https://arxiv.org/pdf/2402.10198>
- ★ Code: <https://github.com/romilbert/samformer>



This project was led by [Romain Ilbert](#) and [myself](#) with our co-authors [Vasilii Feofanov](#), [Aladin Virmaux](#), [Giuseppe Paolo](#), [Themis Palpanas](#), and [Ievgen Redko](#).

## MANO: Exploiting Matrix Norm for Unsupervised Accuracy Estimation Under Distribution Shifts



<https://arxiv.org/pdf/2405.18979>

Thanks for your attention !



- [1] Chen, S.-A., Li, C.-L., Arik, S. O., Yoder, N. C., and Pfister, T. (2023). TSMixer: An all-MLP architecture for time series forecasting. *Transactions on Machine Learning Research*.
- [2] Chen, X., Hsieh, C.-J., and Gong, B. (2022). When vision transformers outperform resnets without pre-training or strong data augmentations. In *International Conference on Learning Representations*.
- [3] Foret, P., Kleiner, A., Mobahi, H., and Neyshabur, B. (2021). Sharpness-aware minimization for efficiently improving generalization. In *International Conference on Learning Representations*.
- [4] Kim, T., Kim, J., Tae, Y., Park, C., Choi, J.-H., and Choo, J. (2021). Reversible instance normalization for accurate time-series forecasting against distribution shift. In *International Conference on Learning Representations*.



- [5] Li, S., Jin, X., Xuan, Y., Zhou, X., Chen, W., Wang, Y.-X., and Yan, X. (2019). Enhancing the locality and breaking the memory bottleneck of transformer on time series forecasting. In Wallach, H., Larochelle, H., Beygelzimer, A., d'Alché-Buc, F., Fox, E., and Garnett, R., editors, *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc.
- [6] Liu, S., Yu, H., Liao, C., Li, J., Lin, W., Liu, A. X., and Dustdar, S. (2022). Pyraformer: Low-complexity pyramidal attention for long-range time series modeling and forecasting. In *International Conference on Learning Representations*.
- [7] Liu, Y., Hu, T., Zhang, H., Wu, H., Wang, S., Ma, L., and Long, M. (2024). itransformer: Inverted transformers are effective for time series forecasting. In *The Twelfth International Conference on Learning Representations*.



- [8] Nie, Y., Nguyen, N. H., Sinthong, P., and Kalagnanam, J. (2023). A time series is worth 64 words: Long-term forecasting with transformers. In *The Eleventh International Conference on Learning Representations*.
- [9] Woo, G., Liu, C., Kumar, A., Xiong, C., Savarese, S., and Sahoo, D. (2024). Unified training of universal time series forecasting transformers.
- [10] Wu, H., Xu, J., Wang, J., and Long, M. (2021). Autoformer: Decomposition transformers with Auto-Correlation for long-term series forecasting. In *Advances in Neural Information Processing Systems*.
- [11] Zeng, A., Chen, M., Zhang, L., and Xu, Q. (2023). Are transformers effective for time series forecasting? In *Proceedings of the AAAI Conference on Artificial Intelligence*.





- [12] Zhai, S., Likhomanenko, T., Littwin, E., Busbridge, D., Ramapuram, J., Zhang, Y., Gu, J., and Susskind, J. M. (2023). Stabilizing transformer training by preventing attention entropy collapse. In Krause, A., Brunskill, E., Cho, K., Engelhardt, B., Sabato, S., and Scarlett, J., editors, *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 40770–40803. PMLR.
- [13] Zhou, H., Zhang, S., Peng, J., Zhang, S., Li, J., Xiong, H., and Zhang, W. (2021). Informer: Beyond efficient transformer for long sequence time-series forecasting. In *The Thirty-Fifth AAAI Conference on Artificial Intelligence, AAAI 2021, Virtual Conference*, volume 35, pages 11106–11115. AAAI Press.



- [14] Zhou, T., Ma, Z., Wen, Q., Wang, X., Sun, L., and Jin, R. (2022). FEDformer: Frequency enhanced decomposed transformer for long-term series forecasting. In *Proc. 39th International Conference on Machine Learning (ICML 2022)*.